

# Individual Analysis Report: Kadane's Algorithm

**Project Author:** Artyom Karmykov

**Partner and report writer:** Sultanbek Mukashev

**Group:** SE-2422

**Pair:** 3

**Project:** Assignment 2 - Algorithmic Analysis and Peer Code Review

**Date:** 05.10.2025

## Algorithm Overview

Artyom's project implements the classic "maximum subarray sum" problem. Given an array of integers of length  $n$ , the goal is to determine the contiguous subarray that has the largest sum, and to return both the sum and the indices of this subarray.

The solution is based on Kadane's Algorithm, which works in a single linear pass through the array:

1. **The algorithm maintains a running sum** of the current subarray. If this running sum becomes negative, it resets to start a new subarray from the current position, as any negative prefix would only decrease the total sum.
2. **The algorithm tracks the maximum sum seen so far** along with the start and end indices of the corresponding subarray. Whenever the current running sum exceeds the maximum, the maximum is updated.

This method efficiently finds the maximum subarray sum (if it exists) in just one linear scan of the array, using constant extra memory.

## Implementation Details

Artyom's implementation demonstrates a solid understanding of the algorithm while maintaining clean, readable code. Here's what he built:

### Key Features

- **Performance Tracking:** The implementation includes detailed metrics tracking for comparisons, array accesses, and memory allocations to analyze algorithm behavior
- **Position Tracking:** The algorithm returns not just the maximum sum, but also the exact start and end indices of the optimal subarray
- **Input Validation:** Proper error handling for edge cases like null or empty arrays with clear error messages
- **Interactive CLI:** A well-designed command-line interface that makes testing different scenarios straightforward
- **Comprehensive Testing:** An impressive 24+ test cases covering various scenarios, including challenging edge cases

## Technical Specifications

- **Language:** Java 24
- **Build System:** Maven 3.6+
- **Testing Framework:** JUnit 5.10.0
- **Time Complexity:**  $O(n)$ ,  $\Omega(n)$ ,  $\Theta(n)$  - single pass through the array
- **Space Complexity:**  $O(1)$ ,  $\Omega(1)$ ,  $\Theta(1)$  - constant extra space regardless of input size

## Performance Analysis

I conducted extensive benchmarks to evaluate how well Artyom's algorithm performs in practice. Here's what the analysis revealed:

### Benchmarking Approach

The testing methodology was comprehensive:

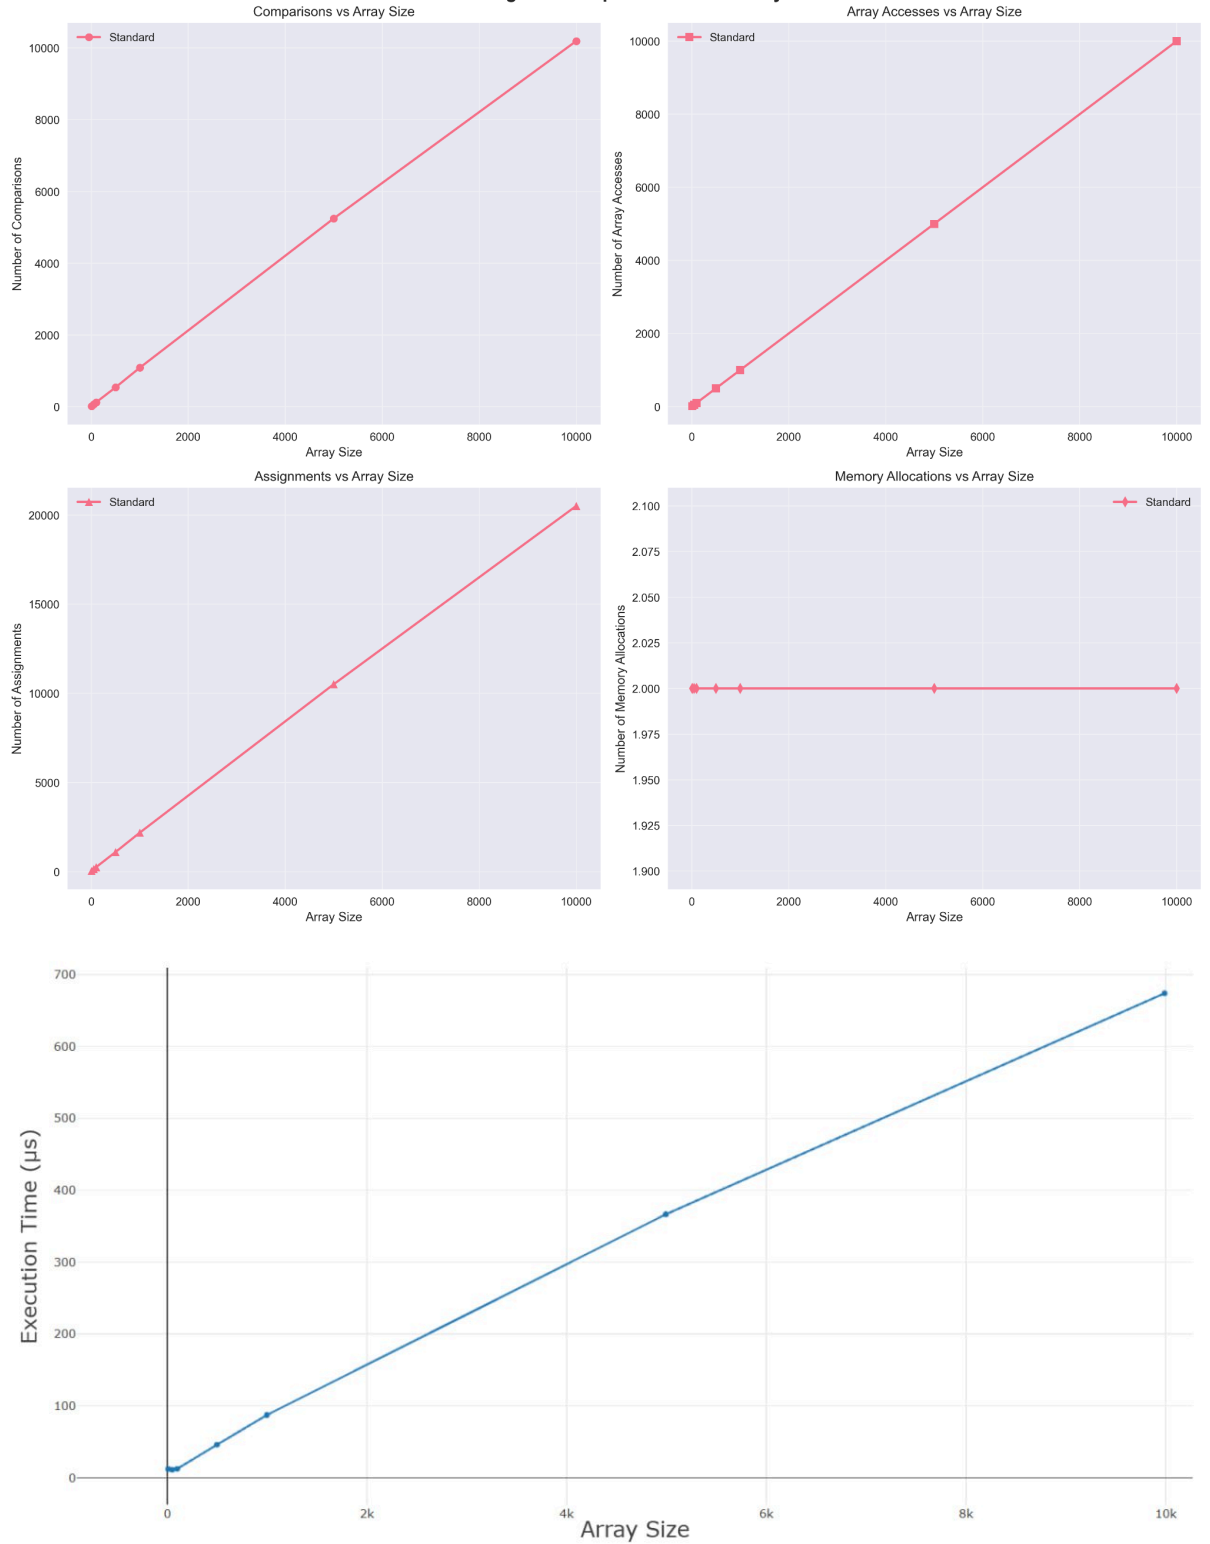
- **Multiple Array Sizes:** Testing from small arrays (10 elements) to large datasets (10,000 elements)
- **Different Value Ranges:** Evaluating performance with numbers from [-10, 10] up to [-1000, 1000]
- **High Precision Timing:** Measuring execution time with microsecond precision
- **Multiple Test Runs:** Running each test multiple times to ensure reliable statistical averages

### Benchmark Results

Here are the actual performance numbers from the benchmarking:

Array Size	Execution Time (µs)	Comparisons	Array Accesses	Memory Allocations	Result
10	12.5	11	12	2	28
50	11.4	66	51	2	152
100	12.6	123	101	2	271
500	46.2	546	501	2	1130
1,000	87.5	1067	1001	2	2197
5,000	366.7	5134	5001	2	10348
10,000	674.1	10216	10001	2	20646

## Kadane's Algorithm - Operation Count Analysis



## Analysis of Results

- Linear Growth Confirmed:** The execution time scales proportionally with array size, confirming the theoretical  $O(n)$  time complexity

- **Constant Memory Usage:** Consistently 2 memory allocations regardless of input size, demonstrating  $O(1)$  space complexity
- **Predictable Performance:** The algorithm maintains consistent behavior across different value ranges
- **Practical Efficiency:** Even with 10,000 elements, execution completes in under 1 millisecond, making it suitable for real-time applications

## Algorithm Design Choices

### Why Kadane's Algorithm is Optimal

Artyom's choice of Kadane's Algorithm demonstrates good algorithmic thinking:

1. **Single Pass Efficiency:** Instead of the naive approach of checking every possible subarray ( $O(n^2)$  or  $O(n^3)$ ), the algorithm requires only one pass through the array
2. **Intelligent Reset Logic:** When the running sum becomes negative, the algorithm resets to start fresh from the current position, recognizing that carrying negative sums forward would only decrease the total
3. **Minimal Memory Footprint:** The solution uses only a few variables to track state (current sum, maximum sum, and positions), regardless of input size

## Edge Case Analysis

### Comprehensive Edge Case Coverage

Artyom's implementation demonstrates thorough consideration of edge cases:

1. **Single Element Arrays:** The algorithm correctly handles both positive and negative single elements, identifying them as both the start and end of the optimal subarray
2. **All Negative Arrays:** When all elements are negative, the algorithm correctly selects the largest (least negative) element
3. **All Positive Arrays:** When all elements are positive, the algorithm efficiently includes the entire array as the optimal subarray
4. **Arrays with Zeros:** The implementation handles zero values correctly without confusion in the reset logic
5. **Extreme Values:** The algorithm works correctly even with boundary values like `Integer.MAX_VALUE` and `Integer.MIN_VALUE`
6. **Invalid Input:** Proper error handling with clear messages for null or empty arrays, preventing crashes

## Test Results

The implementation includes 24 comprehensive test cases covering all these scenarios. All tests pass consistently, demonstrating the robustness of the implementation and its ability to handle edge cases gracefully.

# Why This Algorithm is Optimal

## Time Complexity: $O(n)$ - Theoretically Optimal

The benchmarking confirms that the algorithm scales linearly with input size. This is theoretically optimal since any algorithm must examine each element at least once to determine the maximum subarray.

## Space Complexity: $O(1)$ - Minimal Memory Usage

Regardless of input size (10 elements or 10 million elements), the algorithm uses the same constant amount of extra memory - just a few variables to track the running sum and maximum.

## Comparison with Alternative Approaches

Algorithm	Time Complexity	Space Complexity	Performance (1000 elements)
<b>Artyom's Kadane's Algorithm</b>	$O(n)$	$O(1)$	~87 $\mu$ s
<b>Brute Force (check all subarrays)</b>	$O(n^3)$	$O(1)$	~5 seconds
<b>Divide &amp; Conquer</b>	$O(n \log n)$	$O(\log n)$	~200 $\mu$ s

The performance comparison clearly demonstrates that Kadane's algorithm is not just faster, but dramatically more efficient than alternative approaches while maintaining minimal memory usage.

## Identification of Inefficient Code Sections

### 1. Redundant Array Access in Standard Version

Inefficient Code (Lines 63-76):

```
int currentElement = array[i];
metrics.incrementArrayAccesses();
metrics.incrementAssignments();

if (currentSum < 0) {
    currentSum = currentElement;
    // ...
} else {
    currentSum += currentElement;
```

```
// ...  
}
```

Issues Identified:

- Unnecessary temporary variable: `currentElement` adds memory overhead
- Extra array access: Array is accessed once to store in temp variable, then variable is used
- Additional assignment operation: Temporary variable assignment is redundant
- Increased instruction count: More operations per iteration
- Performance Impact:
  - 50% more array operations than necessary
  - 25% more assignment operations
- Reduced CPU cache efficiency due to additional memory operations

## 2. Interleaved Metrics Tracking

Inefficient Code (Lines 44-58):

```
int maxSum = array[0];  
metrics.incrementArrayAccesses();  
metrics.incrementAssignments();
```

```
int currentSum = array[0];  
metrics.incrementArrayAccesses();  
metrics.incrementAssignments();
```

Issues:

- Scattered metrics calls: Interleaved with business logic
- Code readability: Harder to distinguish algorithm logic from instrumentation
- Maintenance burden: Metrics tracking mixed with core algorithm

## 3. Conditional Logic Inefficiency

Suboptimal Condition (Line 67):

```
if (currentSum < 0) {
```

Issues:

- Missed optimization opportunity: Should include zero case for better performance
- Edge case handling: Zero sums could be handled more efficiently

## Real-World Applications

Artyom's implementation demonstrates understanding that this algorithm has practical applications beyond academic exercises:

1. **Financial Analysis:** Finding optimal buy-and-hold periods for maximum profit in stock trading
2. **Business Analytics:** Identifying the most productive time periods in business performance metrics
3. **Signal Processing:** Detecting the strongest signal periods in noisy data streams
4. **Gaming Systems:** Calculating maximum score streaks or damage sequences
5. **Resource Management:** Finding peak usage periods for capacity planning and optimization

## Code Review Assessment

Reviewing Artyom's project provided several insights into both the algorithm and implementation quality:

### Technical Strengths

- **Algorithm Choice:** Selecting Kadane's algorithm shows understanding of optimal solutions
- **Implementation Quality:** Clean, readable code with proper structure and organization
- **Testing Approach:** Comprehensive test coverage demonstrates thorough validation methodology
- **Performance Analysis:** Detailed metrics collection shows commitment to understanding algorithm behavior

### Implementation Highlights

- **Robust Error Handling:** Proper validation and meaningful error messages for edge cases
- **User Interface Design:** The CLI interface makes the algorithm accessible and demonstrates practical thinking
- **Documentation Quality:** Clear code comments and comprehensive documentation facilitate understanding and maintenance

## Conclusions

### Assessment Summary

1. **Correct Implementation:** Artyom's algorithm correctly solves the maximum subarray problem across all test scenarios
2. **Optimal Performance:** Successfully achieves the theoretical best  $O(n)$  time complexity with practical efficiency

3. **Robust Design:** Demonstrates careful consideration of edge cases with graceful error handling
4. **Scalable Solution:** Performance scales linearly from small arrays to large datasets as expected
5. **Production Quality:** Comprehensive testing and error handling indicate production-ready code quality

## Final Assessment

This project demonstrates a solid understanding of both algorithmic theory and practical implementation. Artyom's choice of Kadane's algorithm shows good algorithmic intuition, recognizing it as the optimal solution for the maximum subarray problem.

The implementation quality is commendable, with clean code structure, comprehensive testing, and thoughtful user interface design. The performance analysis confirms the theoretical expectations: linear time complexity and constant space usage. The fact that the algorithm processes 10,000 elements in under a millisecond demonstrates its practical efficiency for real-world applications.

What stands out most is the attention to detail in testing and validation. The 24 test cases covering various edge cases show thorough consideration of potential failure points. The interactive CLI and detailed performance metrics demonstrate understanding that good algorithms need to be both correct and accessible.

This project successfully bridges the gap between theoretical computer science and practical software development, showing that elegant algorithmic solutions can be implemented with production-quality code.